

✔ Congratulations! You passed!

Grade
received 100%

Latest Submission
Grade 100%

To pass 80% or
higher

Go to next item

1. Why does sequence make a large difference when determining semantics of language?

1 / 1 point

- ☐ It doesn't
- ☐ Because the order of words doesn't matter
- ☒ Because the order in which words appear dictate their impact on the meaning of the sentence
- ☐ Because the order in which words appear dictate their meaning

✔ **Correct**
Correct!

2. How do Recurrent Neural Networks help you understand the impact of sequence on meaning?

1 / 1 point

- ☒ They carry meaning from one cell to the next
- ☐ They don't
- ☐ They look at the whole sentence at a time
- ☐ They shuffle the words evenly

✔ **Correct**
That's right!

3. How does an LSTM help understand meaning when words that qualify each other aren't necessarily beside each other in a sentence?

1 / 1 point

- ☒ Values from earlier words can be carried to later ones via a cell state
- ☐ They load all words into a cell state
- ☐ They shuffle the words randomly
- ☐ They don't

✔ **Correct**
Correct!

4. What keras layer type allows LSTMs to look forward and backward in a sentence?

1 / 1 point

- ☐ Bothdirection
- ☐ Unilateral
- ☒ Bidirectional
- ☐ Bilateral

✔ **Correct**
Correct!

5. What's the output shape of a bidirectional LSTM layer with 64 units?

1 / 1 point


- ☐ (128,1)
- ☐ (128,None)
- ☒ (None, 128)
- ☐ (None, 64)

✔ **Correct**
That's right!

1 / 1 point

6. When stacking LSTMs, how do you instruct an LSTM to feed the next one in the sequence?


- ☒ Ensure that return_sequences is set to True only on units that feed to another LSTM
- ☐ Ensure that they have the same number of units
- ☐ Do nothing, TensorFlow handles this automatically
- ☐ Ensure that return_sequences is set to True on all units

 **Correct**
Correct!

1 / 1 point

7. If a sentence has 120 tokens in it, and a Conv1D with 128 filters with a Kernal size of 5 is passed over it, what's the output shape?

- ☐ (None, 120, 124)
- ☐ (None, 116, 124)
- ☐ (None, 120, 128)
- ☒ (None, 116, 128)

 **Correct**
That's right!

1 / 1 point

8. What's the best way to avoid overfitting in NLP datasets?

- ☐ Use LSTMs
- ☐ Use GRUs
- ☐ Use Conv1D
- ☒ None of the above

 **Correct**
Correct!